GETAMAZEDNOW AI
PRACTICAL AI. POSSIBLE FUTURES.
A DIVISION OF GETAMAZEDNOW PTY LTD

ADSRA · OPEN QUESTION OQ-3

Autonomous sovereignty: who audits the auditor?

The question

At the highest maturity level, AI agents assess and remediate the organisation's own sovereignty posture. What evidence would convince a regulator that this agentic remediation is itself well-governed?

Why it matters

This is the model's most novel contribution and its biggest exposure: the assurance system is itself an AI system, so it needs the same governance as any other AI workload — registered models, scoped tool permissions, budget limits, evidence trails. Without a regulator-credible evidence answer, autonomous assurance stays a slide, not a control. APRA's April 2026 AI letter points the same way: it expects second-line functions able to assess autonomous workflows, and monitoring proportionate to criticality.

Our current position *[a considered view, open to challenge]*

Treat it as aspirational by design — pilots, not production, are the realistic target inside a 24-month planning window. Agentic assurance must be governed as an AI workload, and depends on strong observability as its telemetry substrate.

> **What we're seeking from you**
>
> A concrete evidence package a regulator would accept for autonomous remediation: audit-trail schema, human-oversight thresholds proportionate to blast radius, fail-safe behaviours, and model / tool governance.

Where this is most relevant: Risk & regulatory practitioners · Banking

© 2026 Getamazednow Pty Ltd. All rights reserved